

[NAME](#) | [SYNOPSIS](#) | [DESCRIPTION](#) | [RETURN VALUE](#) | [ERRORS](#) | [CONFORMING TO](#) | [NOTES](#) | [SEE ALSO](#) | [COLOPHON](#)

Search online pages**FLOCK(2)****Linux Programmer's Manual****FLOCK(2)****NAME** [top](#)

flock - apply or remove an advisory lock on an open file

SYNOPSIS [top](#)

```
#include <sys/file.h>

int flock(int fd, int operation);
```

DESCRIPTION [top](#)

Apply or remove an advisory lock on the open file specified by *fd*. The argument *operation* is one of the following:

- LOCK_SH** Place a shared lock. More than one process may hold a shared lock for a given file at a given time.
- LOCK_EX** Place an exclusive lock. Only one process may hold an exclusive lock for a given file at a given time.
- LOCK_UN** Remove an existing lock held by this process.

A call to **flock()** may block if an incompatible lock is held by another process. To make a nonblocking request, include **LOCK_NB** (by ORing) with any of the above operations.

A single file may not simultaneously have both shared and exclusive locks.

Locks created by **flock()** are associated with an open file description (see [open\(2\)](#)). This means that duplicate file descriptors (created by, for example, [fork\(2\)](#) or [dup\(2\)](#)) refer to the same lock, and this lock may be modified or released using any of these file descriptors. Furthermore, the lock is released either by an explicit **LOCK_UN** operation on any of these duplicate file descriptors, or when all such file descriptors have been closed.

If a process uses [open\(2\)](#) (or similar) to obtain more than one file descriptor for the same file, these file descriptors are treated independently by **flock()**. An attempt to lock the file using one of these file descriptors may be denied by a lock that the calling process has already placed via another file descriptor.

A process may hold only one type of lock (shared or exclusive) on a file. Subsequent **flock()** calls on an already locked file will convert an existing lock to the new lock mode.

Locks created by **flock()** are preserved across an **execve(2)**.

A shared or exclusive lock can be placed on a file regardless of the mode in which the file was opened.

RETURN VALUE [top](#)

On success, zero is returned. On error, -1 is returned, and **errno** is set appropriately.

ERRORS [top](#)

EBADF *fd* is not an open file descriptor.

EINTR While waiting to acquire a lock, the call was interrupted by delivery of a signal caught by a handler; see **signal(7)**.

EINVAL *operation* is invalid.

ENOLCK The kernel ran out of memory for allocating lock records.

EWOLDBLOCK

The file is locked and the **LOCK_NB** flag was selected.

CONFORMING TO [top](#)

4.4BSD (the **flock()** call first appeared in 4.2BSD). A version of **flock()**, possibly implemented in terms of **fcntl(2)**, appears on most UNIX systems.

NOTES [top](#)

Since kernel 2.0, **flock()** is implemented as a system call in its own right rather than being emulated in the GNU C library as a call to **fcntl(2)**. With this implementation, there is no interaction between the types of lock placed by **flock()** and **fcntl(2)**, and **flock()** does not detect deadlock. (Note, however, that on some systems, such as the modern BSDs, **flock()** and **fcntl(2)** locks *do* interact with one another.)

In Linux kernels up to 2.6.11, **flock()** does not lock files over NFS (i.e., the scope of locks was limited to the local system). Instead, one could use **fcntl(2)** byte-range locking, which does work over NFS, given a sufficiently recent version of Linux and a server which supports locking. Since Linux 2.6.12, NFS clients support **flock()** ²

locks by emulating them as byte-range locks on the entire file. This means that `fcntl(2)` and `flock()` locks *do* interact with one another over NFS. Since Linux 2.6.37, the kernel supports a compatibility mode that allows `flock()` locks (and also `fcntl(2)` byte region locks) to be treated as local; see the discussion of the `local_lock` option in `nfs(5)`.

`flock()` places advisory locks only; given suitable permissions on a file, a process is free to ignore the use of `flock()` and perform I/O on the file.

`flock()` and `fcntl(2)` locks have different semantics with respect to forked processes and `dup(2)`. On systems that implement `flock()` using `fcntl(2)`, the semantics of `flock()` will be different from those described in this manual page.

Converting a lock (shared to exclusive, or vice versa) is not guaranteed to be atomic: the existing lock is first removed, and then a new lock is established. Between these two steps, a pending lock request by another process may be granted, with the result that the conversion either blocks, or fails if `LOCK_NB` was specified. (This is the original BSD behavior, and occurs on many other implementations.)

SEE ALSO [top](#)

`flock(1)`, `close(2)`, `dup(2)`, `execve(2)`, `fcntl(2)`, `fork(2)`, `open(2)`, `lockf(3)`, `lslocks(8)`

Documentation/filesystems/locks.txt in the Linux kernel source tree (*Documentation/locks.txt* in older kernels)

COLOPHON [top](#)

This page is part of release 4.08 of the Linux *man-pages* project. A description of the project, information about reporting bugs, and the latest version of this page, can be found at <https://www.kernel.org/doc/man-pages/>.

Linux

2016-03-15

FLOCK(2)

[Copyright and license for this manual page](#)

HTML rendering created 2016-10-08 by [Michael Kerrisk](#), author of *The Linux Programming Interface*, maintainer of the *Linux man-pages* project.

For details of in-depth **Linux/UNIX system programming training courses** that I teach, look [here](#).

Hosting by [jambit GmbH](#).

